

```

import requests
from io import StringIO
import pandas as pd
from scipy.io import arff

from scipy.io import arff
import pandas as pd
import requests
from io import StringIO

# URL of the dataset in ARFF format
url = 'http://promise.site.uottawa.ca/SERepository/datasets/jm1.arff'

# Download the dataset
response = requests.get(url)

# Load the dataset from the response content
data = arff.loadarff(StringIO(response.content.decode('utf-8')))
df = pd.DataFrame(data[0])

# Print the first few rows of the DataFrame
print(df.head())

```

	loc	v(g)	ev(g)	iv(g)	n	v	l	d	i	e \
0	1.1	1.4	1.4	1.4	1.3	1.30	1.30	1.30	1.30	1.30
1	1.0	1.0	1.0	1.0	1.0	1.00	1.00	1.00	1.00	1.00
2	72.0	7.0	1.0	6.0	198.0	1134.13	0.05	20.31	55.85	23029.10
3	190.0	3.0	1.0	3.0	600.0	4348.76	0.06	17.06	254.87	74202.67
4	37.0	4.0	1.0	4.0	126.0	599.12	0.06	17.19	34.86	10297.30

	...	l0Code	l0Comment	l0Blank	locCodeAndComment	uniq_Op	uniq_Opnd \
0	...	2.0	2.0	2.0		2.0	1.2
1	...	1.0	1.0	1.0		1.0	1.0
2	...	51.0	10.0	8.0		1.0	17.0
3	...	129.0	29.0	28.0		2.0	17.0
4	...	28.0	1.0	6.0		0.0	11.0

	total_Op	total_Opnd	branchCount	defects
0	1.2	1.2	1.4	b'false'
1	1.0	1.0	1.0	b'true'
2	112.0	86.0	13.0	b'true'
3	329.0	271.0	5.0	b'true'
4	76.0	50.0	7.0	b'true'

[5 rows x 22 columns]

```
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Encode the 'defects' column
df['defects'] = label_encoder.fit_transform(df['defects'])

# Separate the 'defects' column from the rest of the data
defects_column = df['defects']

# Drop the 'defects' column from the DataFrame
df = df.drop(columns=['defects'])

# Initialize StandardScaler
scaler = StandardScaler()

# Fit scaler to your data and transform it
scaled_data = scaler.fit_transform(df)

# Concatenate the scaled features with the 'defects' column
scaled_df = pd.DataFrame(scaled_data, columns=df.columns)
scaled_df['defects'] = defects_column

# Print the first few rows of the scaled DataFrame
print(scaled_df.head())

```

	loc	v(g)	ev(g)	iv(g)	n	v	l \
0	-0.534225	-0.380102	-0.295508	-0.285373	-0.453283	-0.346848	7.255076
1	-0.535530	-0.410826	-0.354578	-0.329250	-0.454485	-0.347003	5.386279
2	0.391486	0.050035	-0.354578	0.219208	0.335124	0.237456	-0.531578
3	1.932161	-0.257206	-0.354578	-0.109867	1.946407	1.895536	-0.469285
4	-0.065494	-0.180396	-0.354578	-0.000175	0.046536	-0.038498	-0.469285

	d	i	e ...	l0Code	l0Comment	l0Blank \
--	---	---	-------	--------	-----------	-----------

```

0 -0.688290 -0.817612 -0.084805 ... -0.406860 -0.081873 -0.263405
1 -0.704325 -0.826329 -0.084806 ... -0.423636 -0.192883 -0.363730
2 0.327797 0.767373 -0.031788 ... 0.415171 0.806207 0.338540
3 0.154084 6.550025 0.086029 ... 1.723710 2.915398 2.345027
4 0.161032 0.157495 -0.061101 ... 0.029320 -0.192883 0.137892

```

```

locCodeAndComment  uniq_Op  uniq_Opnd  total_Op  total_Opnd  branchCount \
0      0.853939 -0.993310 -0.583195 -0.441634 -0.450326 -0.437779
1      0.329798 -1.013221 -0.590695 -0.442954 -0.452319 -0.455481
2      0.329798 0.579644 0.721806 0.289686 0.394739 0.075573
3      0.853939 0.579644 4.434307 1.721964 2.238338 -0.278463
4      -0.194344 -0.017680 -0.028195 0.052073 0.035985 -0.189954

```

```

defects
0      0
1      1
2      1
3      1
4      1

```

[5 rows x 22 columns]

```

print(df.shape) # Check the dimensions of the DataFrame
print(df.head()) # Display the first few rows of the DataFrame

```

```

(10885, 21)
loc  v(g)  ev(g)  iv(g)  n  v  l  d  i  e \
0  1.1  1.4  1.4  1.4  1.3  1.30  1.30  1.30  1.30  1.30
1  1.0  1.0  1.0  1.0  1.0  1.00  1.00  1.00  1.00  1.00
2  72.0  7.0  1.0  6.0  198.0  1134.13  0.05  20.31  55.85  23029.10
3  190.0  3.0  1.0  3.0  600.0  4348.76  0.06  17.06  254.87  74202.67
4  37.0  4.0  1.0  4.0  126.0  599.12  0.06  17.19  34.86  10297.30

...      t  l0Code  l0Comment  l0Blank  locCodeAndComment  uniq_Op \
0  ...  1.30  2.0  2.0  2.0  2.0  1.2
1  ...  1.00  1.0  1.0  1.0  1.0  1.0
2  ...  1279.39  51.0  10.0  8.0  1.0  17.0
3  ...  4122.37  129.0  29.0  28.0  2.0  17.0
4  ...  572.07  28.0  1.0  6.0  0.0  11.0

uniq_Opnd  total_Op  total_Opnd  branchCount
0      1.2      1.2      1.2      1.4
1      1.0      1.0      1.0      1.0
2      36.0     112.0     86.0     13.0
3     135.0     329.0    271.0      5.0
4      16.0      76.0     50.0      7.0

```

[5 rows x 21 columns]

```

# Remove rows with null values
df.dropna(inplace=True)

```

```

# Check for NaN values in each column
nan_counts = df.isna().sum()

```

```

# Display columns with NaN values, if any
columns_with_nan = nan_counts[nan_counts > 0].index.tolist()
if columns_with_nan:
    print("Columns with NaN values:", columns_with_nan)
else:
    print("No NaN values remaining in the DataFrame")

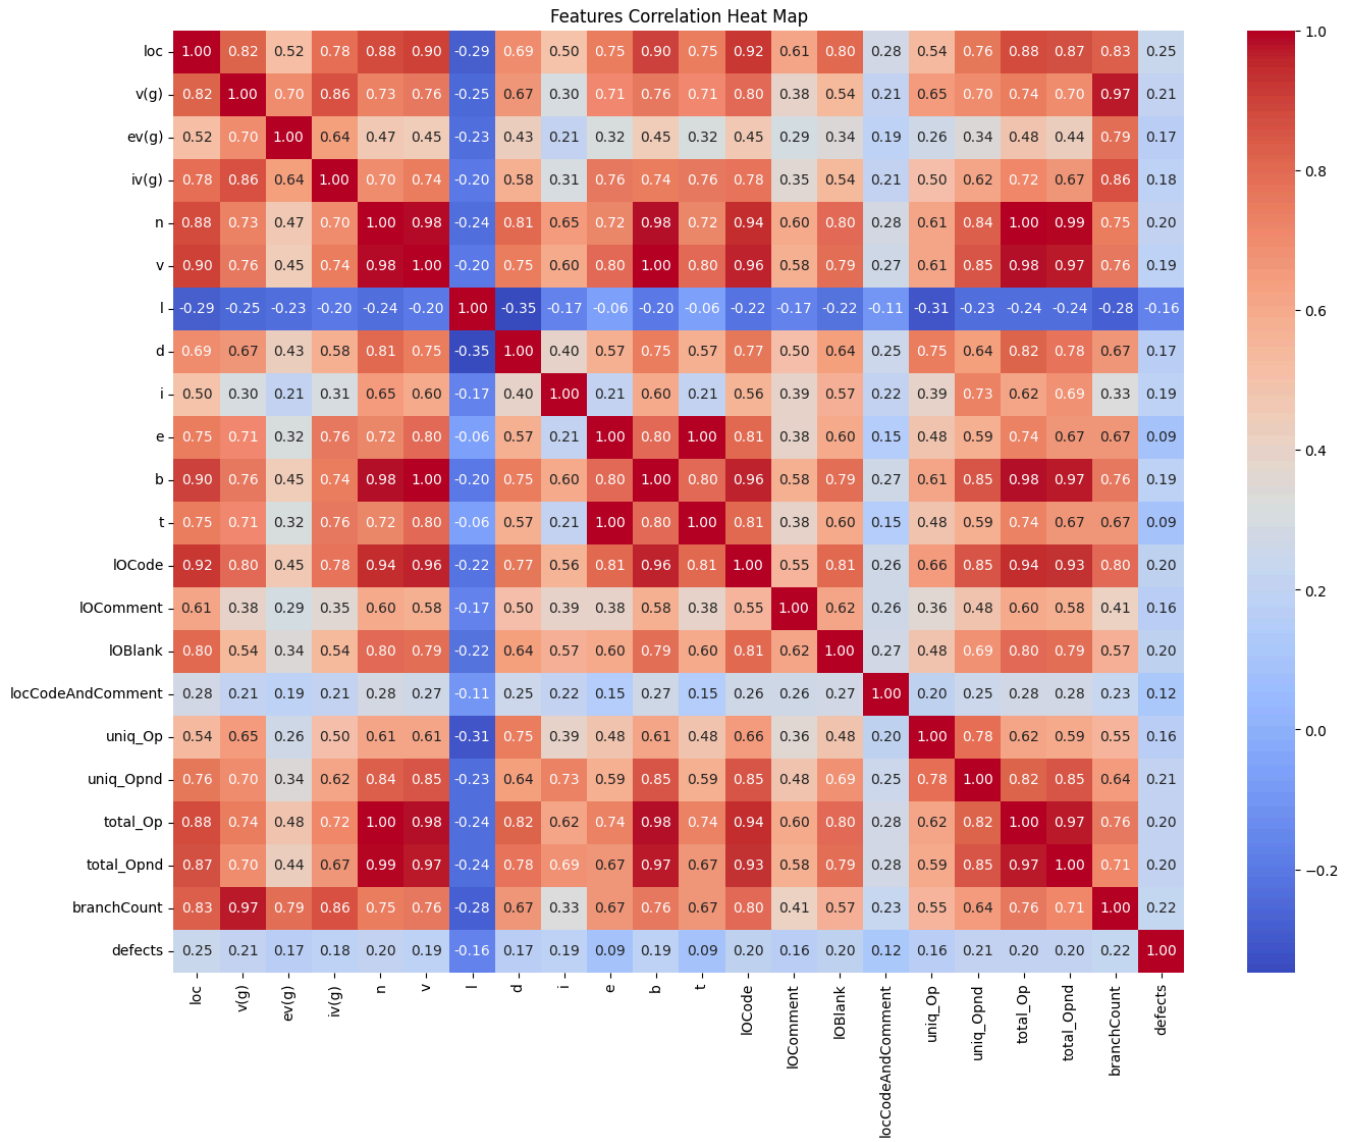
```

No NaN values remaining in the DataFrame

```

# Plotting the features correlation heat map
import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(16, 12))
correlation_matrix = scaled_df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Features Correlation Heat Map')
plt.show()

```



```
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
import matplotlib.pyplot as plt
# Drop the 'defects' column for clustering
clustering_data = scaled_df.drop(columns=['defects'])
# Instantiate the imputer
imputer = SimpleImputer(strategy='mean') # You can choose another strategy if needed

# Impute missing values
clustering_data_imputed = imputer.fit_transform(clustering_data)

# Determine the optimal number of clusters using the elbow method
inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(clustering_data_imputed)
    inertia.append(kmeans.inertia_)

# Plot the elbow curve
plt.plot(range(1, 11), inertia, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method for Optimal K')
plt.show()

# From the elbow curve, choose the optimal number of clusters
optimal_k = 2 # Example: Based on the plot, choose the optimal number of clusters

# Apply K-means clustering with the optimal number of clusters
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
clusters = kmeans.fit_predict(clustering_data_imputed)


# Add cluster labels to the DataFrame
scaled_df['cluster'] = clusters

# Print the first few rows of the DataFrame with cluster labels
print(scaled_df.head())
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(

```



```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 1 in the future. You should set `n_init` to the new value to avoid this warning.
warnings.warn(

```

	loc	v(g)	ev(g)	iv(g)	n	v	l	\
0	-0.534225	-0.380102	-0.295508	-0.285373	-0.453283	-0.346848	7.255076	
1	-0.535530	-0.410826	-0.354578	-0.329250	-0.454485	-0.347003	5.386279	
2	0.391486	0.050035	-0.354578	0.219208	0.335124	0.237456	-0.531578	
3	1.932161	-0.257206	-0.354578	-0.109867	1.946407	1.895536	-0.469285	
4	-0.065494	-0.180396	-0.354578	-0.000175	0.046536	-0.038498	-0.469285	

	d	i	e	...	l0Comment	l0Blank	locCodeAndComment	\
0	-0.688290	-0.817612	-0.084805	...	-0.081873	-0.263405	0.853939	
1	-0.704325	-0.826329	-0.084806	...	-0.192883	-0.363730	0.329798	

```
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier, StackingClassifier
from sklearn.metrics import accuracy_score

# Split the data into training and testing sets
X = scaled_df.drop(columns=['defects', 'cluster']) # Features
y = scaled_df['defects'] # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)

# Initialize classifiers
svm_classifier = SVC(kernel='linear', random_state=42)
nb_classifier = GaussianNB()
rf_classifier = RandomForestClassifier(n_estimators=1000, random_state=42)

# Initialize stacking classifier
estimators = [('svm', svm_classifier), ('nb', nb_classifier), ('rf', rf_classifier)]
stacking_classifier = StackingClassifier(estimators=estimators, final_estimator=nb_classifier)

from sklearn.impute import SimpleImputer

# Instantiate the imputer
imputer = SimpleImputer(strategy='mean') # You can choose another strategy if needed

# Impute missing values in X_train
X_train_imputed = imputer.fit_transform(X_train)

# Impute missing values in X_test
X_test_imputed = imputer.transform(X_test)

# Train the classifiers with imputed data
svm_classifier.fit(X_train_imputed, y_train)
nb_classifier.fit(X_train_imputed, y_train)
rf_classifier.fit(X_train_imputed, y_train)
stacking_classifier.fit(X_train_imputed, y_train)

# Predictions
svm_pred = svm_classifier.predict(X_test_imputed)
nb_pred = nb_classifier.predict(X_test_imputed)
rf_pred = rf_classifier.predict(X_test_imputed)
stacking_pred = stacking_classifier.predict(X_test_imputed)

# Evaluate accuracy
svm_accuracy = accuracy_score(y_test, svm_pred)
nb_accuracy = accuracy_score(y_test, nb_pred)
rf_accuracy = accuracy_score(y_test, rf_pred)
stacking_accuracy = accuracy_score(y_test, stacking_pred)

# Print accuracies
print("SVM Accuracy:", svm_accuracy)
print("Naive Bayes Accuracy:", nb_accuracy)
print("Random Forest Accuracy:", rf_accuracy)
print("Stacking Classifier Accuracy:", stacking_accuracy)
```

```
🔗 SVM Accuracy: 0.8035068188143613
Naive Bayes Accuracy: 0.8054550514890064
Random Forest Accuracy: 0.812134706373504
Stacking Classifier Accuracy: 0.8143612580016699
```

```
from sklearn.metrics import classification_report

# Calculate metrics for SVM
svm_report = classification_report(y_test, svm_pred)
print("SVM Metrics:")
print(svm_report)

# Calculate metrics for Naive Bayes
nb_report = classification_report(y_test, nb_pred)
print("\nNaive Bayes Metrics:")
print(nb_report)

# Calculate metrics for Random Forest
rf_report = classification_report(y_test, rf_pred)
print("\nRandom Forest Metrics:")
print(rf_report)

# Calculate metrics for Stacking Classifier
stacking_report = classification_report(y_test, stacking_pred)
print("\nStacking Classifier Metrics:")
print(stacking_report)
```

↗

SVM Metrics:					
	precision	recall	f1-score	support	
0	0.81	1.00	0.89	2887	
1	0.50	0.01	0.03	706	
accuracy			0.80	3593	
macro avg	0.65	0.51	0.46	3593	
weighted avg	0.75	0.80	0.72	3593	
Naive Bayes Metrics:					
	precision	recall	f1-score	support	
0	0.83	0.95	0.89	2887	
1	0.51	0.21	0.29	706	
accuracy			0.81	3593	
macro avg	0.67	0.58	0.59	3593	
weighted avg	0.77	0.81	0.77	3593	
Random Forest Metrics:					
	precision	recall	f1-score	support	
0	0.83	0.96	0.89	2887	
1	0.56	0.22	0.31	706	
accuracy			0.81	3593	
macro avg	0.69	0.59	0.60	3593	
weighted avg	0.78	0.81	0.78	3593	
Stacking Classifier Metrics:					
	precision	recall	f1-score	support	
0	0.84	0.95	0.89	2887	
1	0.55	0.28	0.37	706	
accuracy			0.81	3593	
macro avg	0.70	0.61	0.63	3593	
weighted avg	0.79	0.81	0.79	3593	

```

import matplotlib.pyplot as plt

# Extracting metrics for SVM
svm_metrics = classification_report(y_test, svm_pred, output_dict=True)
svm_accuracy = svm_metrics['accuracy']
svm_precision = svm_metrics['weighted avg']['precision']
svm_recall = svm_metrics['weighted avg']['recall']
svm_f1_score = svm_metrics['weighted avg']['f1-score']

# Extracting metrics for Naive Bayes
nb_metrics = classification_report(y_test, nb_pred, output_dict=True)
nb_accuracy = nb_metrics['accuracy']
nb_precision = nb_metrics['weighted avg']['precision']
nb_recall = nb_metrics['weighted avg']['recall']
nb_f1_score = nb_metrics['weighted avg']['f1-score']

# Extracting metrics for Random Forest
rf_metrics = classification_report(y_test, rf_pred, output_dict=True)
rf_accuracy = rf_metrics['accuracy']
rf_precision = rf_metrics['weighted avg']['precision']
rf_recall = rf_metrics['weighted avg']['recall']
rf_f1_score = rf_metrics['weighted avg']['f1-score']

# Extracting metrics for Stacking Classifier
stacking_metrics = classification_report(y_test, stacking_pred, output_dict=True)
stacking_accuracy = stacking_metrics['accuracy']
stacking_precision = stacking_metrics['weighted avg']['precision']
stacking_recall = stacking_metrics['weighted avg']['recall']
stacking_f1_score = stacking_metrics['weighted avg']['f1-score']

# Plotting
labels = ['Accuracy', 'Precision', 'Recall', 'F1-Score']
svm_values = [svm_accuracy, svm_precision, svm_recall, svm_f1_score]
nb_values = [nb_accuracy, nb_precision, nb_recall, nb_f1_score]
rf_values = [rf_accuracy, rf_precision, rf_recall, rf_f1_score]
stacking_values = [stacking_accuracy, stacking_precision, stacking_recall, stacking_f1_score]

x = range(len(labels))


plt.figure(figsize=(12, 8)) # Increase the figure size

plt.bar(x, svm_values, width=0.2, label='SVM', align='center')
plt.bar([i + 0.2 for i in x], nb_values, width=0.2, label='Naive Bayes', align='center')
plt.bar([i + 0.4 for i in x], rf_values, width=0.2, label='Random Forest', align='center')
plt.bar([i + 0.6 for i in x], stacking_values, width=0.2, label='Stacking Classifier', align='center')

plt.xlabel('Metrics')
plt.ylabel('Score')
plt.title('Comparison of Classification Metrics for Different Algorithms')
plt.xticks([i + 0.3 for i in x], labels)
plt.legend(loc='upper left', bbox_to_anchor=(1, 1)) # Move legend outside the plot

plt.tight_layout() # Adjust subplots to fit into figure area.
plt.show()

```

prompt: i want to calculate MAE, RMSE, and MAPE, of applied algorithms

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_absolute_percentage_error
```

```
# Calculate mean absolute error (MAE)
svm_mae = mean_absolute_error(y_test, svm_pred)
nb_mae = mean_absolute_error(y_test, nb_pred)
rf_mae = mean_absolute_error(y_test, rf_pred)
stacking_mae = mean_absolute_error(y_test, stacking_pred)
```

```
# Calculate root mean squared error (RMSE)
svm_rmse = mean_squared_error(y_test, svm_pred, squared=False)
nb_rmse = mean_squared_error(y_test, nb_pred, squared=False)
rf_rmse = mean_squared_error(y_test, rf_pred, squared=False)
stacking_rmse = mean_squared_error(y_test, stacking_pred, squared=False)
```

```
# Print the calculated metrics
print("SVM MAE:", svm_mae)
print("Naive Bayes MAE:", nb_mae)
print("Random Forest MAE:", rf_mae)
print("Stacking Classifier MAE:", stacking_mae)

print("\nSVM RMSE:", svm_rmse)
print("Naive Bayes RMSE:", nb_rmse)
print("Random Forest RMSE:", rf_rmse)
print("Stacking Classifier RMSE:", stacking_rmse)
```



```
SVM MAE: 0.19649318118563874
Naive Bayes MAE: 0.1945449485109936
Random Forest MAE: 0.18786529362649595
Stacking Classifier MAE: 0.18563874199833008

SVM RMSE: 0.4432755138575091
Naive Bayes RMSE: 0.4410724980215765
Random Forest RMSE: 0.4334343013958355
```

Stacking Classifier RMSE: 0.4308581460275877

```
import matplotlib.pyplot as plt
import numpy as np

# Define the algorithms and their corresponding metrics
algorithms = ['SVM', 'Naive Bayes', 'Random Forest', 'Stacking Classifier']
```