

```

import requests
from io import StringIO
import pandas as pd
from scipy.io import arff

from scipy.io import arff
import pandas as pd
import requests
from io import StringIO

# URL of the dataset in ARFF format
url = 'http://promise.site.uottawa.ca/SERepository/datasets/kc1.arff'

# Download the dataset
response = requests.get(url)

# Load the dataset from the response content
data = arff.loadarff(StringIO(response.content.decode('utf-8')))
df = pd.DataFrame(data[0])

# Print the first few rows of the DataFrame
print(df.head())

```

	loc	v(g)	ev(g)	iv(g)	n	v	l	d	i	e	...	\
0	1.1	1.4	1.4	1.4	1.3	1.30	1.30	1.30	1.30	1.30	...	...
1	1.0	1.0	1.0	1.0	1.0	1.00	1.00	1.00	1.00	1.00	...	...
2	83.0	11.0	1.0	11.0	171.0	927.89	0.04	23.04	40.27	21378.61	...	...
3	46.0	8.0	6.0	8.0	141.0	769.78	0.07	14.86	51.81	11436.73	...	...
4	25.0	3.0	1.0	3.0	58.0	254.75	0.11	9.35	27.25	2381.95	...	...

	locCode	locComment	locBlank	locCodeAndComment	uniq_Op	uniq_Opnd	\
0	2.0	2.0	2.0		2.0	1.2	1.2
1	1.0	1.0	1.0		1.0	1.0	1.0
2	65.0	10.0	6.0		0.0	18.0	25.0
3	37.0	2.0	5.0		0.0	16.0	28.0
4	21.0	0.0	2.0		0.0	11.0	10.0

	total_Op	total_Opnd	branchCount	defects
0	1.2	1.2	1.4	b'false'
1	1.0	1.0	1.0	b'true'
2	107.0	64.0	21.0	b'true'
3	89.0	52.0	15.0	b'true'
4	41.0	17.0	5.0	b'true'

[5 rows x 22 columns]

```

from sklearn.preprocessing import StandardScaler, LabelEncoder

```

```

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Encode the 'defects' column
df['defects'] = label_encoder.fit_transform(df['defects'])

# Separate the 'defects' column from the rest of the data
defects_column = df['defects']

# Drop the 'defects' column from the DataFrame
df = df.drop(columns=['defects'])

# Initialize StandardScaler
scaler = StandardScaler()

# Fit scaler to your data and transform it
scaled_data = scaler.fit_transform(df)

# Concatenate the scaled features with the 'defects' column
scaled_df = pd.DataFrame(scaled_data, columns=df.columns)
scaled_df['defects'] = defects_column

# Print the first few rows of the scaled DataFrame
print(scaled_df.head())

```

	loc	v(g)	ev(g)	iv(g)	n	v	l	\
0	-0.647864	-0.368740	-0.124739	-0.339674	-0.580634	-0.498642	3.093249	
1	-0.651226	-0.471309	-0.306546	-0.458191	-0.584224	-0.499223	2.146739	
2	2.105319	2.092900	-0.306546	2.504721	1.449754	1.296396	-0.882093	
3	0.861512	1.323638	1.966040	1.615847	1.090817	0.990097	-0.787442	
4	0.155568	0.041533	-0.306546	0.134392	0.097757	-0.007646	-0.661241	

	d	i	e	...	locCode	locComment	locBlank	\
--	---	---	---	-----	---------	------------	----------	---

```

0 -0.695929 -0.927649 -0.300506 ... -0.517950 0.341722 0.062345
1 -0.734088 -0.941606 -0.300524 ... -0.559302 0.017524 -0.196995
2 2.069348 0.885308 0.925197 ... 2.087232 2.935302 1.099707
3 1.028871 1.422170 0.355163 ... 0.929373 0.341722 0.840366
4 0.328012 0.279593 -0.164008 ... 0.267740 -0.306673 0.062345

```

```

locCodeAndComment  uniq_Op  uniq_Opnd  total_Op  total_Opnd  branchCount  \
0      2.652867 -1.122654 -0.683788 -0.576537 -0.548440 -0.419224
1      1.232121 -1.157565 -0.700191 -0.580400 -0.554677 -0.470570
2     -0.188624 1.809801 1.268178 1.467364 1.409972 2.096706
3     -0.188624 1.460699 1.514224 1.119630 1.035753 1.326523
4     -0.188624 0.587944 0.037947 0.192341 -0.055719 0.042885

```

```

defects
0      0
1      1
2      1
3      1
4      1

```

```
[5 rows x 22 columns]
```

```

print(df.shape) # Check the dimensions of the DataFrame
print(df.head()) # Display the first few rows of the DataFrame

```

```

(2109, 21)
loc  v(g)  ev(g)  iv(g)  n  v  l  d  i  e ... \
0  1.1  1.4  1.4  1.4  1.3  1.30  1.30  1.30  1.30  1.30 ...
1  1.0  1.0  1.0  1.0  1.0  1.00  1.00  1.00  1.00  1.00 ...
2  83.0  11.0  1.0  11.0  171.0  927.89  0.04  23.04  40.27  21378.61 ...
3  46.0  8.0  6.0  8.0  141.0  769.78  0.07  14.86  51.81  11436.73 ...
4  25.0  3.0  1.0  3.0  58.0  254.75  0.11  9.35  27.25  2381.95 ...

```

```

t  l0Code  l0Comment  l0Blank  locCodeAndComment  uniq_Op  uniq_Opnd  \
0  1.30  2.0  2.0  2.0  2.0  1.2  1.2
1  1.00  1.0  1.0  1.0  1.0  1.0  1.0
2  1187.70  65.0  10.0  6.0  0.0  18.0  25.0
3  635.37  37.0  2.0  5.0  0.0  16.0  28.0
4  132.33  21.0  0.0  2.0  0.0  11.0  10.0

```

```

total_Op  total_Opnd  branchCount
0      1.2      1.2      1.4
1      1.0      1.0      1.0
2     107.0     64.0     21.0
3      89.0     52.0     15.0
4      41.0     17.0      5.0

```

```
[5 rows x 21 columns]
```

```

# Remove rows with null values
df.dropna(inplace=True)

```

```

# Check for NaN values in each column
nan_counts = df.isna().sum()

```

```

# Display columns with NaN values, if any
columns_with_nan = nan_counts[nan_counts > 0].index.tolist()
if columns_with_nan:
    print("Columns with NaN values:", columns_with_nan)
else:
    print("No NaN values remaining in the DataFrame")


```

```
No NaN values remaining in the DataFrame
```

```

# Plotting the features correlation heat map
import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(16, 12))
correlation_matrix = scaled_df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Features Correlation Heat Map')
plt.show()

```



```
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Drop the 'defects' column for clustering
clustering_data = scaled_df.drop(columns=['defects'])

# Determine the optimal number of clusters using the elbow method
inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(clustering_data)
    inertia.append(kmeans.inertia_)

# Plot the elbow curve
plt.plot(range(1, 11), inertia, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method for Optimal K')
plt.show()

# From the elbow curve, choose the optimal number of clusters
optimal_k = 2 # Example: Based on the plot, choose the optimal number of clusters

# Apply K-means clustering with the optimal number of clusters
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
clusters = kmeans.fit_predict(clustering_data)


# Add cluster labels to the DataFrame
scaled_df['cluster'] = clusters

# Print the first few rows of the DataFrame with cluster labels
print(scaled_df.head())
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(

```



```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarni
warnings.warn(
loc      v(g)      ev(g)      iv(g)      n      v      l  \
0 -0.647864 -0.368740 -0.124739 -0.339674 -0.580634 -0.498642  3.093249
1 -0.651226 -0.471309 -0.306546 -0.458191 -0.584224 -0.499223  2.146739
2  2.105319  2.092900 -0.306546  2.504721  1.449754  1.296396 -0.882093
3  0.861512  1.323638  1.966040  1.615847  1.090817  0.990097 -0.787442
4  0.155568  0.041533 -0.306546  0.134392  0.097757 -0.007646 -0.661241

d      i      e      ...  10Comment  10Blank  locCodeAndComment  \

```

```

from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier, StackingClassifier
from sklearn.metrics import accuracy_score

# Split the data into training and testing sets
X = scaled_df.drop(columns=['defects', 'cluster']) # Features
y = scaled_df['defects'] # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)

# Initialize classifiers
svm_classifier = SVC(kernel='linear', random_state=42)
nb_classifier = GaussianNB()
rf_classifier = RandomForestClassifier(n_estimators=1000, random_state=42)

# Initialize stacking classifier
estimators = [('svm', svm_classifier), ('nb', nb_classifier), ('rf', rf_classifier)]
stacking_classifier = StackingClassifier(estimators=estimators, final_estimator=nb_classifier)

# Train the classifiers
svm_classifier.fit(X_train, y_train)
nb_classifier.fit(X_train, y_train)
rf_classifier.fit(X_train, y_train)
stacking_classifier.fit(X_train, y_train)

# Predictions
svm_pred = svm_classifier.predict(X_test)
nb_pred = nb_classifier.predict(X_test)
rf_pred = rf_classifier.predict(X_test)
stacking_pred = stacking_classifier.predict(X_test)

# Evaluate accuracy
svm_accuracy = accuracy_score(y_test, svm_pred)
nb_accuracy = accuracy_score(y_test, nb_pred)
rf_accuracy = accuracy_score(y_test, rf_pred)
stacking_accuracy = accuracy_score(y_test, stacking_pred)

# Print accuracies
print("SVM Accuracy:", svm_accuracy)
print("Naive Bayes Accuracy:", nb_accuracy)
print("Random Forest Accuracy:", rf_accuracy)
print("Stacking Classifier Accuracy:", stacking_accuracy)

```

```

SVM Accuracy: 0.8591954022988506
Naive Bayes Accuracy: 0.8074712643678161
Random Forest Accuracy: 0.8548850574712644
Stacking Classifier Accuracy: 0.8160919540229885

```

```

from sklearn.metrics import classification_report

# Calculate metrics for SVM
svm_report = classification_report(y_test, svm_pred)
print("SVM Metrics:")
print(svm_report)

# Calculate metrics for Naive Bayes
nb_report = classification_report(y_test, nb_pred)
print("\nNaive Bayes Metrics:")
print(nb_report)

# Calculate metrics for Random Forest
rf_report = classification_report(y_test, rf_pred)
print("\nRandom Forest Metrics:")
print(rf_report)

# Calculate metrics for Stacking Classifier
stacking_report = classification_report(y_test, stacking_pred)
print("\nStacking Classifier Metrics:")
print(stacking_report)

```

```

SVM Metrics:

```

	precision	recall	f1-score	support
0	0.86	0.99	0.92	588
1	0.71	0.16	0.26	108
accuracy			0.86	696
macro avg	0.79	0.57	0.59	696
weighted avg	0.84	0.86	0.82	696

Naive Bayes Metrics:				
	precision	recall	f1-score	support
0	0.88	0.90	0.89	588
1	0.36	0.32	0.34	108
accuracy			0.81	696
macro avg	0.62	0.61	0.62	696
weighted avg	0.80	0.81	0.80	696

Random Forest Metrics:				
	precision	recall	f1-score	support
0	0.88	0.96	0.92	588
1	0.56	0.30	0.39	108
accuracy			0.85	696
macro avg	0.72	0.63	0.65	696
weighted avg	0.83	0.85	0.84	696

Stacking Classifier Metrics:				
	precision	recall	f1-score	support
0	0.89	0.90	0.89	588
1	0.40	0.38	0.39	108
accuracy			0.82	696
macro avg	0.64	0.64	0.64	696
weighted avg	0.81	0.82	0.81	696

```

import matplotlib.pyplot as plt

# Extracting metrics for SVM
svm_metrics = classification_report(y_test, svm_pred, output_dict=True)
svm_accuracy = svm_metrics['accuracy']
svm_precision = svm_metrics['weighted avg']['precision']
svm_recall = svm_metrics['weighted avg']['recall']
svm_f1_score = svm_metrics['weighted avg']['f1-score']

# Extracting metrics for Naive Bayes
nb_metrics = classification_report(y_test, nb_pred, output_dict=True)
nb_accuracy = nb_metrics['accuracy']
nb_precision = nb_metrics['weighted avg']['precision']
nb_recall = nb_metrics['weighted avg']['recall']
nb_f1_score = nb_metrics['weighted avg']['f1-score']

# Extracting metrics for Random Forest
rf_metrics = classification_report(y_test, rf_pred, output_dict=True)
rf_accuracy = rf_metrics['accuracy']
rf_precision = rf_metrics['weighted avg']['precision']
rf_recall = rf_metrics['weighted avg']['recall']
rf_f1_score = rf_metrics['weighted avg']['f1-score']

# Extracting metrics for Stacking Classifier
stacking_metrics = classification_report(y_test, stacking_pred, output_dict=True)
stacking_accuracy = stacking_metrics['accuracy']
stacking_precision = stacking_metrics['weighted avg']['precision']
stacking_recall = stacking_metrics['weighted avg']['recall']
stacking_f1_score = stacking_metrics['weighted avg']['f1-score']

# Plotting
labels = ['Accuracy', 'Precision', 'Recall', 'F1-Score']
svm_values = [svm_accuracy, svm_precision, svm_recall, svm_f1_score]
nb_values = [nb_accuracy, nb_precision, nb_recall, nb_f1_score]
rf_values = [rf_accuracy, rf_precision, rf_recall, rf_f1_score]
stacking_values = [stacking_accuracy, stacking_precision, stacking_recall, stacking_f1_score]

x = range(len(labels))

plt.figure(figsize=(12, 8)) # Increase the figure size

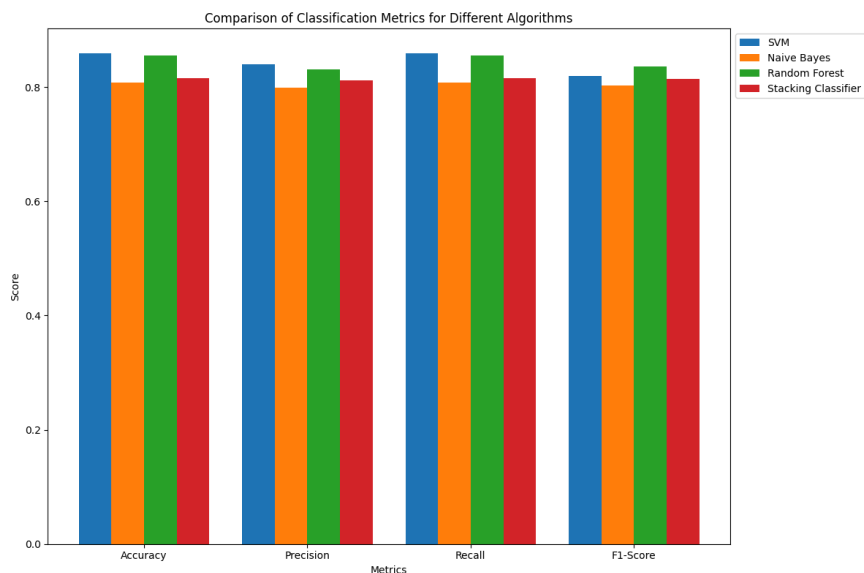
plt.bar(x, svm_values, width=0.2, label='SVM', align='center')
plt.bar([i + 0.2 for i in x], nb_values, width=0.2, label='Naive Bayes', align='center')
plt.bar([i + 0.4 for i in x], rf_values, width=0.2, label='Random Forest', align='center')
plt.bar([i + 0.6 for i in x], stacking_values, width=0.2, label='Stacking Classifier', align='center')

plt.xlabel('Metrics')
plt.ylabel('Score')
plt.title('Comparison of Classification Metrics for Different Algorithms')
plt.xticks([i + 0.3 for i in x], labels)
plt.legend(loc='upper left', bbox_to_anchor=(1, 1)) # Move legend outside the plot

plt.tight_layout() # Adjust subplots to fit into figure area.
plt.show()

```





# prompt: i want to calculate MAE, RMSE, and MAPE, of applied algorithms

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_absolute_percentage_error
```

```
# Calculate mean absolute error (MAE)
svm_mae = mean_absolute_error(y_test, svm_pred)
nb_mae = mean_absolute_error(y_test, nb_pred)
rf_mae = mean_absolute_error(y_test, rf_pred)
stacking_mae = mean_absolute_error(y_test, stacking_pred)

# Calculate root mean squared error (RMSE)
svm_rmse = mean_squared_error(y_test, svm_pred, squared=False)
nb_rmse = mean_squared_error(y_test, nb_pred, squared=False)
rf_rmse = mean_squared_error(y_test, rf_pred, squared=False)
stacking_rmse = mean_squared_error(y_test, stacking_pred, squared=False)
```

```
# Print the calculated metrics
print("SVM MAE:", svm_mae)
print("Naive Bayes MAE:", nb_mae)
print("Random Forest MAE:", rf_mae)
print("Stacking Classifier MAE:", stacking_mae)

print("\nSVM RMSE:", svm_rmse)
print("Naive Bayes RMSE:", nb_rmse)
print("Random Forest RMSE:", rf_rmse)
print("Stacking Classifier RMSE:", stacking_rmse)
```



```
SVM MAE: 0.14080459770114942
Naive Bayes MAE: 0.1925287356321839
Random Forest MAE: 0.14511494252873564
Stacking Classifier MAE: 0.1839080459770115

SVM RMSE: 0.3752393871932282
Naive Bayes RMSE: 0.4387809654396871
Random Forest RMSE: 0.38093955232915316
Stacking Classifier RMSE: 0.4288450139351179
```